Package 'beast'

November 29, 2017

Type Package

Title Bayesian Estimation of Change-Points in the Slope of
Multivariate Time-Series

Version 1.0

Date 2017-11-29

Author Panagiotis Papastamoulis

Maintainer Panagiotis Papastamoulis <papapast@yahoo.gr>

Description Assume that a temporal process is composed of contiguous segments with differ-
ing slopes and replicated noise-corrupted time series measurements are observed. The un-
known mean of the data generating process is modelled as a piecewise linear func-
tion of time with an unknown number of change-points. The package infers the joint poste-
rior distribution of the number and position of change-points as well as the unknown mean pa-
rameters per time-series by MCMC sampling. A-priori, the proposed model uses an overfit-
ting number of mean parameters but, conditionally on a set of change-points, only a sub-
set of them influences the likelihood. An exponentially decreasing prior distribution on the num-
ber of change-points gives rise to a posterior distribution concentrating on sparse representa-
tions of the underlying sequence, but also available is the Poisson distribution. See Papasta-
moulis et al (2017) <arXiv:1709.06111> for a detailed presentation of the method.

License GPL-2

Imports RColorBrewer

Depends R (>= 2.10)

NeedsCompilation no

R topics documented:

| beast-package | *Bayesian Estimation of Change-Points in the Slope of Multivariate Time-Series* |

Description

Assume that a temporal process is composed of contiguous segments with differing slopes and replicated noise-corrupted time series measurements are observed. The unknown mean of the data generating process is modelled as a piecewise linear function of time with an unknown number of change-points. The package infers the joint posterior distribution of the number and position of change-points as well as the unknown mean parameters per time-series by MCMC sampling. A-priori, the proposed model uses an overfitting number of mean parameters but, conditionally on a set of change-points, only a subset of them influences the likelihood. An exponentially decreasing prior distribution on the number of change-points gives rise to a posterior distribution concentrating on sparse representations of the underlying sequence, but also available is the Poisson distribution. See Papastamoulis et al (2017) <arXiv:1709.06111> for a detailed presentation of the method.

Details

The beast package deals with **B**ayesian **e**stimation of ch**a**nge-points in the **s**lope of multivariate **t**ime-series, introduced by Papastamoulis et al (2017). For a given period $t = 1, \ldots, T$ we observe multiple time series which are assumed independent, each one consisting of multiple measurements (replicates). Each time series is assumed to have its own segmentation, which is common among its replicates. Thus, different time series have distinct mean parameters in the underlying normal distribution. The variance, which is assumed known, can be either shared between different time series or not and in practice it is estimated at a pre-processing stage.

Our model infers the joint posterior distribution of the number and location of change-points by MCMC sampling. For this purpose a Metropolis-Hastings MCMC sampler is used. The main function of the package is beast.

Assume that the observed data consists of N time-series, each one consisting of R variables (or replicates) measured at T consecutive time-points. The input of the main function [beast](#) should be given in the form of a list myDataList with the following attributes:

- length(myDataList) should be equal to R, that is, the number of variables (or replicates)
- dim(myDataList)[[1]],..., dim(myDataList)[[R]] should be all $T \times N$, that is, rows and columns should correspond to time-points and different series, respectively.

Then, a basic usage of the package consists of the following commands:

- beastRun <- beast(myDataList = myDataList)
- print(beastRun)
- plot(beastRun)

which correspond to running the MCMC sampler, printing and plotting output summaries, respectively.

Author(s)

Panagiotis Papastamoulis

Maintainer: Panagiotis Papastamoulis <papapast@yahoo.gr>

References

Papastamoulis P., Furukawa T., van Rhijn N., Bromley M., Bignell E. and Rattray M. (2017). Bayesian detection of piecewise linear trends in replicated time-series with application to growth data modelling. arXiv:**1709.06111** [stat.AP]

See Also

[beast](#), [print.beast.object](#), [plot.beast.object](#)

Examples

```
# toy-example (MCMC iterations not enough)
library('beast') # load package
data("FungalGrowthDataset") # load dataset
myIndex <- c(392, 62, 3, 117) # run the sampler only for the
#                                 specific subset of time-series
set.seed(1)
# Run MCMC sampler with very small number of iterations (nIter):
run_mcmc <- beast(myDataList = FungalGrowthDataset, subsetIndex = myIndex,
zeroNormalization = TRUE, nIter = 40, burn = 20)
# Print output:
print(run_mcmc)
# Plot output to file: "beast_plot.pdf"
plot(run_mcmc, fileName = "beast_plot_toy.pdf", timeScale=1/6, xlab = "hours", ylab = "growth")
# Run the following commands to obtain convergence:

## Not run:
# This example illustrates the package using a subset of four
```

```
#       time-series of the fungal dataset.
library('beast') # load package
data("FungalGrowthDataset") # load dataset
myIndex <- c(392, 62, 3, 117) # run the sampler only for the
#                                  specific subset of time-series
set.seed(1) # optional
# Run MCMC sampler with the default number of iterations (nIter =70000):
run_mcmc <- beast(myDataList = FungalGrowthDataset, subsetIndex = myIndex,
zeroNormalization = TRUE)
# Print output:
print(run_mcmc)
# Plot output to file: "beast_plot.pdf"
plot(run_mcmc, fileName = "beast_plot.pdf", timeScale=1/6, xlab = "hours", ylab = "growth")
# NOTE 1: for a complete analysis remove the `subsetIndex = myIndex` argument.
# NOTE 2: `zeroNormalization = TRUE` is an optional argument that forces all
#    time-series to start from zero. It is not supposed to be used
#    for other applications.

## End(Not run)
```

---

beast   *Main function*

Description

This is the main function of the package, implementing the MCMC sampler described in Papastamoulis et al (2017).

Usage

```
beast(myDataList, burn, nIter, mhPropRange, mhSinglePropRange, startPoint,
    timeScale, savePlots, zeroNormalization, LRange, tau,
    gammaParameter, nu0, alpha0, beta0, subsetIndex, saveTheta, sameVariance,
    Prior
)
```

Arguments

myDataList	Observed data in the form of a list with length R, denoting the dimensionality of the multivariate time-series data. For each $r = 1, \ldots, R$, myDataList[[r]] should correspond to $T \times N$ array, with myDataList[[r]][t, n] corresponding to the observed data for time-series $n = 1, \ldots, N$ and time-point $t = 1, \ldots, T$.
burn	Number of iterations that will be discarder as burn-in period. Default value: burn = 20000.
nIter	Number of MCMC iterations. Default value: nIter = 70000.
mhPropRange	Positive integer corresponding to the parameter d_1 of MCMC Move 3.a of Papastamoulis et al (2017). Default value: mhPropRange = 1.

mhSinglePropRange

Positive integer denoting the parameter d_2 of Papastamoulis et al (2017). Default value: `mhPropRange = 40`.

startPoint   An (optional) positive integer pointing at the minimum time-point where changes are allowed to occur. Default value: `startPoint = 2` (all possible values are taken into account).

timeScale   Null.

savePlots   Character denoting the name of the folder where various plots will be saved to.

zeroNormalization

Logical value denoting whether to normalize to zero all time time-series for $t = 1$. Default: `zeroNormalization = FALSE`.

LRange   Range of possible values for the number of change-points. Default value: `LRange = 0:30`.

tau   Positive real number corresponding to parameter c in Move 2 of Papastamoulis et al (2017). Default value: `tau = 0.05`.

gammaParameter   Positive real number corresponding to parameter α of the exponential prior distribution. Default value: `gammaParameter = 2`.

nu0   Positive real number corresponding to prior parameter ν_0 in Papastamoulis et al (2017). Default value: `nu0 = 0.1`.

alpha0   Positive real number corresponding to prior parameter α_0 in Papastamoulis et al (2017). Default value: `alpha0 = 1`.

beta0   Positive real number corresponding to prior parameter β_0 in Papastamoulis et al (2017). Default value: `beta0 = 1`.

subsetIndex   Optional subset of integers corresponding to time-series indexes. If not null, the sampler will be applied only to the specified subset.

saveTheta   Logical value indicating whether to save the MCMC output. Default: True.

sameVariance   Logical value indicating whether to assume the same variance per time-point across time-series. Default value: `sameVariance = TRUE`.

Prior   Character string specifying the prior distribution of the number of change-points. Allowed values: `Prior = "complexity"` (default) or `Prior = "Poisson"` (not suggested).

Value

The output of the sampler is returned as a list, with the following features:

Cutpoint_posterior_median

The estimated medians per change-point, conditionally on the most probable number of change-points per time-series.

Cutpoint_posterior_variance

The estimated variances per change-points, conditionally on the most probable number of change-points per time-series.

NumberOfCutPoints_posterior_distribution

Posterior distributions of number of change-points per time-series.

NumberOfCutPoints_MAP

The most probable number of change-points per time-series.

```
Metropolis-Hastings_acceptance_rate
                Acceptance of the MCMC move-types.
```

subject_ID   the identifier of individual time-series.

```
Cutpoint_mcmc_trace_map
```
                The sampled values of each change-point per time series, conditionally on the
                MAP values.

theta   The sampled values of the means per time-series, conditionally on the MAP
                values.

```
nCutPointsTrace
```
                The sampled values of the number of change-points, per time-series.

Note

The complexity prior distribution with parameter gammaParameter = 2 is the default prior as-
sumption imposed on the number of change-points. Smaller (larger) values of gammaParameter
will a-priori support larger (respectively: smaller) number of change-points.

For completeness purposes, the Poisson distribution is also allowed in the Prior. In this latter case,
the gammaParameter denotes the rate parameter of the Poisson distribution. Note that in this case
the interpretation of gammaParameter is reversed: Smaller (larger) values of gammaParameter will
a-priori support smaller (respectively: larger) number of change-points.

Author(s)

Panagiotis Papastamoulis

References

Papastamoulis P., Furukawa T., van Rhijn N., Bromley M., Bignell E. and Rattray M. (2017).
Bayesian detection of piecewise linear trends in replicated time-series with application to growth
data modelling. arXiv:1709.06111 [stat.AP]

Examples

```
# toy-example (MCMC iterations not enough)
library('beast') # load package
data("FungalGrowthDataset") # load dataset
myIndex <- c(392, 62, 3, 117) # run the sampler only for the
#                                    specific subset of time-series
set.seed(1)
# Run MCMC sampler with very small number of iterations (nIter):
run_mcmc <- beast(myDataList = FungalGrowthDataset, subsetIndex = myIndex,
zeroNormalization = TRUE, nIter = 40, burn = 20)
# Print output:
print(run_mcmc)
# Plot output to file: "beast_plot.pdf"
plot(run_mcmc, fileName = "beast_plot_toy.pdf", timeScale=1/6, xlab = "hours", ylab = "growth")
# Run the following commands to obtain convergence:

## Not run:
# This example illustrates the package using a subset of four
```

```
#       time-series of the fungal dataset.
library('beast') # load package
data("FungalGrowthDataset") # load dataset
myIndex <- c(392, 62, 3, 117) # run the sampler only for the
#                               specific subset of time-series
set.seed(1) # optional
# Run MCMC sampler with the default number of iterations (nIter =70000):
run_mcmc <- beast(myDataList = FungalGrowthDataset, subsetIndex = myIndex,
zeroNormalization = TRUE)
# Print output:
print(run_mcmc)
# Plot output to file: "beast_plot.pdf"
plot(run_mcmc, fileName = "beast_plot.pdf", timeScale=1/6, xlab = "hours", ylab = "growth")
# NOTE 1: for a complete analysis remove the `subsetIndex = myIndex` argument.
# NOTE 2: `zeroNormalization = TRUE` is an optional argument that forces all
#    time-series to start from zero. It is not supposed to be used
#    for other applications.

## End(Not run)
```

---

birthProbs   *Birth Probabilities*

Description

This function defines the probability of proposing an addition of a change-point.

Usage

```
birthProbs(LRange)
```

Arguments

LRange   The range of possible values for the number of change-points.

Value

probs   Vector of probabilities

Author(s)

Panagiotis Papastamoulis

---

complexityPrior   *Complexity prior distribution*

Description

This function computes the complexity prior distribution on the number of change-points, defined as $f(\ell) = P(\ell_n = \ell) \propto e^{-\alpha \ell \log(bT/\ell)}, a, b > 0; \ell = 0, 1, 2, \ldots$. Note that this distribution has exponential decrease (Castillo and van der Vaart, 2012) when $b > 1 + e$, so we set $b = 3$.

Usage

```
complexityPrior(Lmax = 20, gammaParameter, nTime)
```

Arguments

Lmax   maximum number of change-points (default = 20).

gammaParameter positive real number, corresponding to α.

nTime   positive integer denoting the total number of time-points.

Value

logPrior   Prior distribution values in the log-scale.

Author(s)

Panagiotis Papastamoulis

References

Castillo I. and van der Vaart A (2012). Needles and Straw in a Haystack: Posterior concentration for possibly sparse sequences. The Annals of Statistics, 40(4), 2069–2101.

---

computeEmpiricalPriorParameters
                    Compute the empirical mean.

Description

This function computes the empirical mean of the time-series.

Usage

```
computeEmpiricalPriorParameters(myDataList, nu0 = 1, alpha0 = 1, beta0 = 1)
```

Arguments

`myDataList`	Observed multivariate time-series.
`nu0`	positive real number.
`alpha0`	positive real number.
`beta0`	positive real number.

Value

`mu0`	Empirical mean

Author(s)

Panagiotis Papastamoulis

`computePosteriorParameters`

Compute empirical posterior parameters

Description

Compute empirical posterior parameters

Usage

```
computePosteriorParameters(myDataList, priorParameters)
```

Arguments

`myDataList`	Observed data.
`priorParameters`	
	Prior parameters.

Value

list of posterior parameters

Author(s)

Panagiotis Papastamoulis

---

computePosteriorParametersFree

Posterior parameters

Description

Posterior parameters

Usage

computePosteriorParametersFree(myDataList, priorParameters)

Arguments

myDataList   observed data.

priorParameters

list of prior parameters.

Value

list of posterior parameters.

Author(s)

Panagiotis Papastamoulis

FungalGrowthDataset   *Fungal Growth Dataset*

Description

Time-series dataset with $N \times R \times T$ growth levels for $R = 3$ replicates of $N = 411$ objects (mutants) measured every 10 minutes for $T = 289$ time-points. See Papastamoulis et al (2017) for a detailed description.

Usage

FungalGrowthDataset

Format

Time-series data.

References

Papastamoulis P., Furukawa T., van Rhijn N., Bromley M., Bignell E. and Rattray M. (2017). Bayesian detection of piecewise linear trends in replicated time-series with application to growth data modelling. arXiv:**1709.06111** [stat.AP]

`localProposal` *Move 3.b*

Description

Implements Move 3.b of the Metropolis-Hastings MCMC sampler.

Usage

```
localProposal(cutPoints, nTime, mhPropRange, startPoint)
```

Arguments

cutPoints	Current configuration of change-points.
nTime	Total number of time-points.
mhPropRange	Parameter d_2.
startPoint	Integer, at least equal to 2.

Value

newState	Candidate state of the chain.
propRatio	Proposal ratio.

Author(s)

Panagiotis Papastamoulis

`logLikelihoodFullModel`

Log-likelihood function.

Description

Log-likelihood function.

Usage

```
logLikelihoodFullModel(myData, cutPoints, theta, startPoint)
```

Arguments

myData	data
cutPoints	change-points.
theta	means.
startPoint	optional integer at least equal to 2.

Value

log-likelihood value.

Author(s)

Panagiotis Papastamoulis

---

logPrior   *Log-prior.*

Description

Logarithm of the prior distribution on the number of change-points.

Usage

```
logPrior(cutPoints, nTime, startPoint)
```

Arguments

cutPoints	change-points.
nTime	number of time-points.
startPoint	optional integer, at least equal to 2.

Value

logarithm of the prior distribution.

Author(s)

Panagiotis Papastamoulis

---

mcmcSampler   *MCMC sampler*

Description

This function implements the Metropolis-Hastings MCMC sampler for individual time-series.

Usage

```
mcmcSampler(myData, nIter, finalIterationPdf, modelVariance, mhPropRange,
mhSinglePropRange, movesRange, startPoint, postPar, dName, timeScale,
burn, iterPerPlotPrefix, priorParameters, L = 3, LRange, tau,
gammaParameter, saveTheta, Prior = "complexity")
```

Arguments

`myData`	observed data.
`nIter`	number of mcmc iterations
`finalIterationPdf`	
	output folder
`modelVariance`	null
`mhPropRange`	positive integer
`mhSinglePropRange`	
	positive integer
`movesRange`	null
`startPoint`	positive integer
`postPar`	list of emprirically estimated parameters
`dName`	subject ID
`timeScale`	null
`burn`	burn-in period.
`iterPerPlotPrefix`	
	null
`priorParameters`	
	prior parameters.
`L`	null
`LRange`	range of possible values of the number of change-points.
`tau`	real.
`gammaParameter`	real.
`saveTheta`	TRUE.
`Prior`	character.

Value

MCMC output.

Author(s)

Panagiotis Papastamoulis

myUnicodeCharacters   *Printing*

Description

Printing various unicode symbols.

Usage

```
myUnicodeCharacters()
```

Value

printed symbol

normalizeTime0   *Zero normalization*

Description

Zero normalization at 1st time-point

Usage

```
normalizeTime0(myDataList)
```

Arguments

myDataList   data

Value

null

Author(s)

Panagiotis Papastamoulis

---

plot.beast.object *Plot function*

Description

This function plots objects returned by the beast function. All output is diverted to a pdf file provided in the fileName argument.

Usage

```
## S3 method for class 'beast.object'
plot(x, fileName, width, height, pointsize, ylab, xlab, timeScale, myPal, ...)
```

Arguments

x	An object of class beast.object, which is returned by the beast function.
fileName	Name of the output pdf file.
width	Width of pdf file. Default: 9
height	Height pdf file. Default: 6
pointsize	Pointsize. Default: 12
ylab	y-axis label. Default: x.
xlab	x-axis label. Default: t.
timeScale	A multiplicative-factor which will be used to scale the x-axis labels. For example, if time-points correspond to 10-minute periods, then setting timeScale = 1/6 will make the x-axis labels correspond to hours. Default: timeScale = 1.
myPal	Vector of colors that will be used to produce the plot with all time-series overlayed. If the distinct values of the inferred numbers of change-points is less than 10, the Set1 pallete of the RColorBrewer library is used. Otherwise, the user has to manually define the colors.
...	ignored.

Details

The function will produce a plot with all time-series coloured according to the corresponding number of change-points. In addition, it will generate individual plots per time-series displaying the observed data with boxplots which summarize the posterior distribution of change-points locations, conditionally on the most probable number of change-points.

Author(s)

Panagiotis Papastamoulis

---

print.beast.object   *Print function*

Description

This function prints a summary of objects returned by the beast function.

Usage

```
## S3 method for class 'beast.object'
print(x, ...)
```

Arguments

x	An object of class beast.object, which is returned by the beast function.
...	ignored.

Details

The function prints a summary of the most probable number (MAP) of change-points per time-series in the form of a table, as well as a list containing the MAP number of change-points and the corresponding locations (posterior medians) per time-series.

Author(s)

Panagiotis Papastamoulis

---

proposeTheta   *Move 2*

Description

Proposes an update of θ according to Metropolis-Hastings move 2.

Usage

```
proposeTheta(thetaOld, tau, alpha, beta)
```

Arguments

thetaOld	Current values
tau	Parameter c.
alpha	null
beta	null

Value

mean   proposed values.

Author(s)

Panagiotis Papastamoulis

---

simMultiIndNormInvGamma

Prior random numbers

Description

Generation of mean values according to the prior

Usage

```
simMultiIndNormInvGamma(mu, nu, alpha, beta)
```

Arguments

mu   means

nu   precision parameter

alpha   prior parameter

beta   prior parameter

Value

null

Author(s)

Panagiotis Papastamoulis

---

simulateFromPrior   *Generate change-points according to the prior*

Description

Generate change-points according to the prior distribution conditionally on a given number of change-points.

Usage

```
simulateFromPrior(nTime, startPoint, L = 3)
```

Arguments

nTime	Number of time-points
startPoint	At least equal to 2.
L	null

Value

cutPoints	Change-point locations

Author(s)

Panagiotis Papastamoulis

---

singleLocalProposal   *Move 3.b*

Description

Implement Metropolis-Hastings Move 3.b.

Usage

```
singleLocalProposal(cutPoints, nTime, mhSinglePropRange, startPoint)
```

Arguments

cutPoints	Current state
nTime	Number of time-points
mhSinglePropRange	
	Prior parameter.
startPoint	Optional.

Value

newState	Candidate state
propRatio	Proposal ratio

Author(s)

Panagiotis Papastamoulis

---

truncatedPoisson *Truncated Poisson pdf*

Description

Probability density function of the truncated Poisson distribution.

Usage

```
truncatedPoisson(Lmax = 20, gammaParameter = 1)
```

Arguments

Lmax	Max number
gammaParameter	Location parameter.

Value

logPrior	Log-pdf values

Author(s)

Panagiotis Papastamoulis

---

updateNumberOfCutpoints
Move 1

Description

Update the number of change-points according to Metropolis-Hastings move 1.

Usage

```
updateNumberOfCutpoints(cutPoints, nTime, startPoint, LRange, birthProbs)
```

Arguments

cutPoints	Current configuration
nTime	Number of time-points
startPoint	Optional integer
LRange	Range of possible values
birthProbs	Birth probabilities

Value

newState	Candidate state
propRatio	Proposal ratio

Author(s)

Panagiotis Papastamoulis

Index